

# Programming Assignment 1

W. Daniel Hiromoto

10-01-2024

## Contents

<b>Program 1 - Affine Map</b>	<b>2</b>
Image Data Structure . . . . .	2
Transformation Algorithm . . . . .	3
Define size of output image . . . . .	3
Inverse of Transformation Matrix . . . . .	3
Iterate over every pixel in Output image . . . . .	4
Bilinear Interpolation Implementation . . . . .	4

Code can be found at [Github](#)

## Program 1 - Affine Map

We have the following affine map:

$$\begin{bmatrix} x1 \\ x2 \end{bmatrix} \rightarrow \begin{bmatrix} 1 & 1 \\ 1 & 3 \end{bmatrix} \begin{bmatrix} x1 \\ x2 \end{bmatrix} + \begin{bmatrix} 4 \\ 0 \end{bmatrix}$$

This affine map can be converted into homogenous coordinates in 3 Dimensions to include the translation:

$$\begin{bmatrix} 1 & 1 & 4 \\ 1 & 3 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

This can then be decomposed into the following transformation matrices:

$$\text{Shear: } \begin{bmatrix} 1 & 2 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \text{ Rotation: } \begin{bmatrix} \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} & 0 \\ \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} & 0 \\ 0 & 0 & 1 \end{bmatrix} \text{ Scale: } \begin{bmatrix} \sqrt{2} & 0 & 0 \\ 0 & \sqrt{2} & 0 \\ 0 & 0 & 1 \end{bmatrix} \text{ Translation: } \begin{bmatrix} 1 & 0 & 4 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

In order to apply these transformation to an image in python, we'll first import a few external libraries, namely, `numpy`, `imageio`, and `matplotlib.pyplot`

After loading an image of a black square on a white background, we need to iterate over every pixel of this image. Before we can do so, we need to understand how images are stored once read by the imageio library.

```
i1 = iio.imread('blackSquare.png')
```



Figure 1: 50 x 50 Black Square on White Background

## Image Data Structure

Instead of every pixel value having a coordinate, it is placed in an array, with the “Y-Coordinates” Stored first followed by the “X-Coordinates”. This data structure will be a 3D array, as the color data (RGB) of every pixel is stored in an array of length 3. For simplicity sake, I will only be working with images that do not have an alpha/transparency channel:

```
[
    [[255, 255, 255], [255, 255, 255], ..., [255, 255, 255]],
    [[0, 0, 0], [0, 0, 0], ..., [0, 0, 0]],
    [[255, 255, 255], [255, 255, 255], ..., [255, 255, 255]],
]
```

Data Structure of images when read

The “Coordinate” values will correspond to indexes in the image array. I.e. `i1[y, x, :]` will get the pixel data at coordinate x,y.

Finally, the origin begins at the top left of the image, and the positive Y coordinates go to the bottom of the image and positive X coordinates go the right of the image.

## Transformation Algorithm

The algorithm for applying these transformations will go as follow:

1. Define the size of the output image
2. Take the inverse of a given transformation matrix
3. Iterate over every pixel in a defined output image
4. For every pixel in our output image, find it's equivalent location in the original image using the inverse transform
5. Implement the Bilinear Interpolation algorithm to determine the value of the pixel in the output image

### Define size of output image

The is very simple, I added the ability to both enter `height,width` dimensions manually,

```
def imgTransform(src, matrix, outputSize=None):  
    ...  
    # set output size manually if set  
    if outputSize is not None:  
        height = outputSize[0]  
        width = outputSize[1]
```

or I attempt to find the minimum size needed by multiplying the fareset corner of the original image by the provided transform. This does not work in all cases for rotations, but for time's sake this was successful most of the time:

```
def imgTransform(src, matrix, outputSize=None):  
    # calculate smallest possible size needed  
    height, width, _ = src.shape  
    width, height, _ = (matrix @ [width-1, height-1, 1]).astype(int)
```

Once I have the size of the output image, I can initialize it. Here, I set all pixels to 0 (black) for simplicity:

```
# init output values  
output = np.zeros((height, width, 3), dtype=np.uint8)
```

### Inverse of Transformation Matrix

To save on time, numpy has an implentation of this that can do it for us:

```
# take inverse of matrix  
inverse = np.linalg.inv(matrix)
```

## Iterate over every pixel in Output image

I could have iterated over every pixel from the source image, but it's easier to iterate over the output image. If I had iterated over the source image, I would have to determine if the transformed pixel from source to output is within bounds of the output image. Furthermore, if the transform is a skew, magnify, scale, or rotate, you would be left with gaps in the output image. To fill in the gaps, you would have to iterate over the output image to interpolate the pixels. If I have to go back and iterate over the output image to interpolate, there is little reason to iterate over the source image. This would explain why many image processing APIs (such as [OpenCV](#) and [Scipy](#)) will ask for or convert to, the inverse of the transform matrix.

```
# iterate rows
for y, row in enumerate(output):
    # iterate columns
    for x, pixel in enumerate(row):
        # Apply to new image
        output[y, x, :] = bilinearInterpolation(src, inverse, y, x)
return output
```

End of 'imgTransform' function

I use Bilinear Interpolation to determine the value of every pixel in the output image.

## Bilinear Interpolation Implementation

Interpolation aims to solve the following problem:

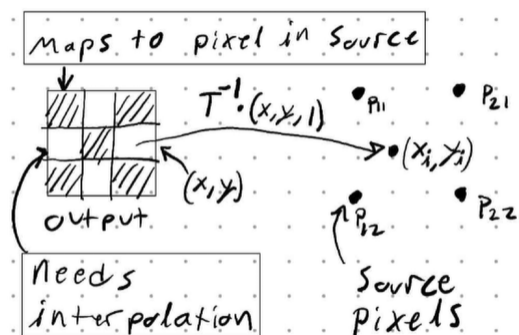


Figure 2: Interpolation problem

Where  $T^{-1}$  is the inverse transform matrix. Unshaded squares with coordinates  $x, y$  are pixels in the output image that do not exist as a pixel in the original image after the inverse transformation is applied. The Points  $P$  in the image are pixels from the original image.

After applying the inverse transformation matrix the resulting source image coordinates,  $x_i, y_i$ , will lie somewhere between the pixels of the source image.

```
def bilinearInterpolation(src, inverse, y, x):
    ...
    # setup the coordinate vector from output image
    coord = np.array([x, y, 1])

    # find the point in the initial image
    xi, yi, _ = (inverse @ coord)
```

In Bilinear Interpolation, the value will effectively be the weighted mean of the surrounding pixel values:




Figure 3: text

```
# setup coordinates of pixels
x1 = np.floor(xi).astype(int)
x2 = np.ceil(xi).astype(int)
y1 = np.floor(yi).astype(int)
y2 = np.ceil(yi).astype(int)
...
# get the original surrounding values
p11 = src[y1, x1, :]
p12 = src[y2, x1, :]
p21 = src[y1, x2, :]
p22 = src[y2, x2, :]
```